

Automatic differentiation in Julia

Miles Lubin and Jarrett Revels

17th Euro AD Workshop

August 20, 2015

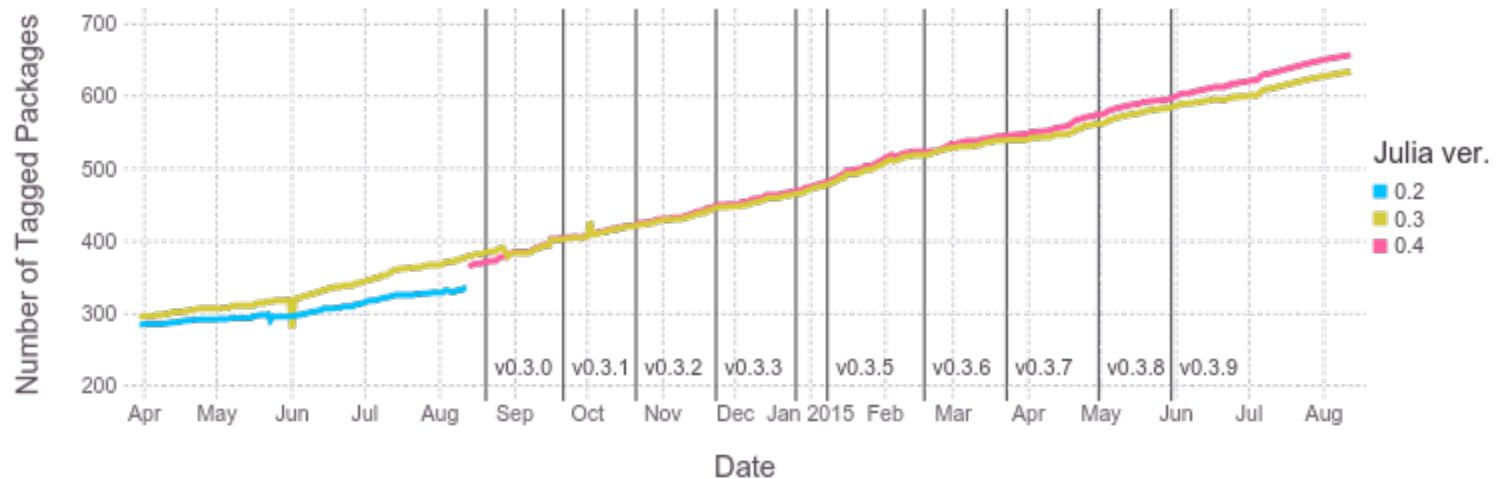
Why Julia?

- Fast like C++, high level like Python and Matlab
 - That's the idea at least
- Solves the “multiple-language problem” in technical computing

Julia timeline

- Julia publicly announced, 2012
 - Julia 0.1 release, February 2013
 - Julia 0.2 release, November 2013
- 1st annual JuliaCon held in Chicago, 2014
 - Julia 0.3 release, August 2014
- 2nd annual JuliaCon held at MIT, 2015
- Julia 0.4 release, Soon™

It's not 1.0, but people find it useful...



Total number of packages by Julia version

JuliaDiff is a:

- [Web page](#)
- [Github organization](#)

to help organize the development of AD tools in Julia.

Most people interact with JuliaDiff through:

```
optimize(f, method=:l_bfgs, autodiff=true)
```

```
function rosenbrock100(x::Vector)
    out = zero(eltype(x))
    for i in 1:div(length(x),2)
        out += 100*(x[2i-1]^2 - x[2i])^2 + (x[2i-1]-1)^2
    end
    out
end
```

```
@time optimize(rosenbrock100, zeros(100),
    method = :l_bfgs, iterations=21)
```

```
# elapsed time: 0.003834211 seconds
```

```
# Value of Function at Minimum: 3.419262
```

```
@time optimize(rosenbrock100, zeros(100),
    method = :l_bfgs, iterations=21, autodiff=true)
```

```
# elapsed time: 0.002318992 seconds
```

```
# Value of Function at Minimum: 0.000000
```

Outline

- Introduction to technical features of Julia interesting for AD
- ForwardDiff package
- JuMP modeling language for optimization

Follow along

<https://juliabox.org/>

<https://github.com/mlubin/EuroAD2015>

JuMP - a modeling language for linear and nonlinear optimization

$$\begin{array}{ll}\min & f(x) \\ \text{s.t.} & g(x) \leq 0 \\ & h(x) = 0\end{array}$$

- *All functions given as closed-form algebraic expressions*

State of the art

Commercial tools:

- AMPL (Gay, Fourer)
 - De-facto standard .nl exchange format
- GAMS

Open source:

- Pyomo
 - Writes to .nl format, doesn't implement AD
- YALMIP
 - Not large scale, no hessians
- CasADi

What JuMP looks like...

```
m = Model(solver=IpoptSolver())  
@defVar(m, x[1:n])  
@setNLObjective(m, Min,  
    sum{ exp(x[i]^2), i = 1:n} )  
@addNLConstraint(m,  
    prod{ x[i], i=1:n} <= 1)  
solve(m)
```

JuMP is a domain-specific language

```
myset = ["cat", "dog"]  
@defVar(m, x[myset])  
@addNLConstraint(m,  
    sin(x["dog"]) <= 0.5)
```

Useful for indexing over:

- Edges in a graph
- Types of widgets to produce
- ...

JuMP is a domain-specific language

```
@defVar(m, 1[i] ≤ x[i=1:N] ≤ 2i)
@defVar(m, t ≥ 0)
@addNLConstraint(m, limit[i=(3:N-1)],
                 exp(x[i]) ≤ t)
```

Equivalent to

```
for i in 3:N-1
    @addNLConstraint(m, exp(x[i]) ≤ t)
end
```

Easy to query derivatives

```
m = Model(); @defVar(m, x); @defVar(m, y)
@setNLObjective(m, Min, sin(x) + sin(y))
```

```
values = [2.0, 3.0]
```

```
d = JuMPNLPEvaluator(m); initialize(d, [:Grad])
```

```
objval = eval_f(d, values)
```

```
 $\nabla f$  = zeros(2)
```

```
eval_grad_f(d,  $\nabla f$ , values)
```

```
#  $\nabla f$  == [cos(2.0), cos(3.0)]
```

Benchmarks

- Build model in memory, prepare AD
 - **Model generation time**
- Give to Ipopt for 5 iterations, report time spent in NLP evaluations (incl. **gradients, jacobians, hessian of the lagrangian**)

<https://github.com/mlubin/JuMPSupplement>

Model generation time (sec.)

Instance	JuMP	Commercial		Open-source	
		AMPL	GAMS	Pyomo	YALMIP
clnlbeam-5	9	0	0	5	117
clnlbeam-50	11	2	3	43	>600
clnlbeam-500	28	21	34	424	>600
acpower-1	22	0	0	3	-
acpower-10	28	1	6	26	-
acpower-100	54	16	471	263	-

Derivative evaluation time (sec.)

Instance	JuMP	Commercial	
		AMPL	GAMS
clnlbeam-5	0.03	0.03	0.09
clnlbeam-50	0.39	0.34	0.74
clnlbeam-500	4.72	3.40	15.69
acpower-1	0.08	0.02	0.19
acpower-10	0.81	0.35	5.07
acpower-100	9.28	3.42	424.89

What do we do?

- `sum{}` and `prod{}` translated to accumulation loops
 - AMPL flattens out
- Apply Reverse-mode AD to this expression graph
 - Recompute instead of storing intermediate terms inside loops
 - Fuse reverse and forward mode of top-level loops (thanks Paul)

What do we do?

- Applying reverse mode, use Julia's code generation facilities to **generate and compile a function at runtime** which evaluates the gradient
- This gives us gradients and Jacobians

From gradients to Hessians

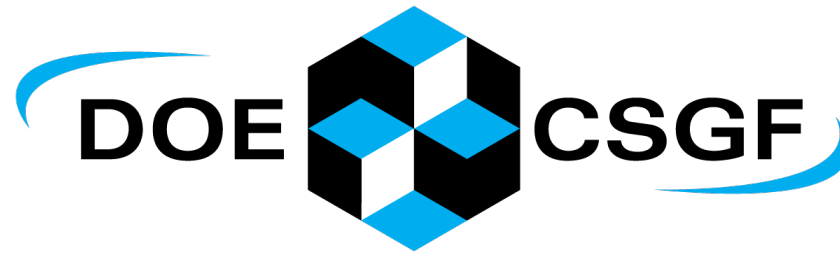
- Apply forward-mode to gradient functions to evaluate Hessian-vector product
- Acyclic graph coloring heuristic of Gebremedhin et. al (2009)

$$\begin{bmatrix} h_{11} & h_{12} & & h_{14} & \\ h_{12} & h_{22} & h_{23} & & \\ & h_{23} & h_{33} & & \\ h_{14} & & & h_{44} & \\ & & & & h_{55} \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} + h_{14} \\ h_{12} + h_{23} & h_{22} \\ h_{33} & h_{23} \\ h_{14} & h_{44} \\ & h_{55} \end{bmatrix}$$

Discussion

- Loops versus flattened out expression graphs
 - Algebraic simplifications? Presolve?
 - Composability

Thanks!



GORDON AND BETTY
MOORE
FOUNDATION